

MLP

🧠 Model: Multi-Layer Perceptron (MLP)

Dataset and Preprocessing

The dataset used corresponds to the Bose–Hubbard model simulations at fixed hopping $t = 1.0$, covering the interaction and chemical potential range $\bar{U} \in [0, 25]$ and $\bar{\mu} \in [-5, 60]$ with a step size of 0.1, resulting in approximately **163,000 points**. Each record contains physical observables including:

Density (ρ), Φ^2 (superfluid order parameter), Energy, and iteration count.

The **phase label** was determined using a physics-based criterion:

- **Mott Insulator (0)**: if $\rho \approx \text{integer}$ and $\Phi^2 < 10^{-6}$,
- **Superfluid (1)**: otherwise.

Thus, the classification task maps the phase of the system directly from the two macroscopic control parameters $(\bar{U}, \bar{\mu})$.

Data Reduction (Sampling Strategy)

While the full dataset contains >160k samples, several **reduced versions** were prepared to study data-efficiency and generalization behavior:

Dataset	Total Samples	Composition
Full	~160k	Complete domain coverage
Reduced (30k)	30,684	Balanced sampling across MI, SF, and boundary
Reduced (10k)	10,684	Moderate down-sampling
Reduced (6k)	6,684	Aggressive sampling — minimal representation

Each reduced dataset was constructed by **selective sampling**:

- 15k (MI) + 15k (SF) + 30k (boundary) for 30k set (balanced regions),
- proportional scaling for smaller sets.

This reduction tests whether the MLP can still reconstruct the **phase boundary** and maintain classification accuracy with fewer physics samples.

Model Architecture

The MLP model was implemented using **scikit-learn's** `MLPClassifier` inside a `Pipeline` with feature standardization:

python

 Copy code

```
Pipeline([
    ("scaler", StandardScaler()),
    ("clf", MLPClassifier(hidden_layer_sizes=(50, 50),
                          max_iter=500, random_state=42))
])
```

Architecture:

- **Input layer:** 2 neurons ($\bar{U}$, $\bar{M}$)
- **Hidden layers:** (50, 50) neurons with ReLU activation
- **Output layer:** 1 neuron with sigmoid activation (binary classification)
- **Optimizer:** Adam (default learning rate 10^{-3})
- **Loss:** Cross-entropy
- **Batch size:** Auto (scikit-learn adaptive)
- **Max iterations:** 500

This small, smooth architecture ensures regularization and avoids overfitting while being expressive enough to capture nonlinear phase boundaries.

Training and Evaluation

The model was trained with an **80–20 train-test split** under **stratified sampling**, ensuring both MI and SF regions were proportionally represented.

Evaluation was conducted at three levels:

1. **On test set** within sampled data
2. **Extrapolation test** — predictions on a fine grid extended beyond training region ($\bar{U} \leq 30, \bar{\mu} \leq 65$)
3. **Full-dataset comparison** — matching predicted grid with full simulation data

Results

(a) Full Dataset (~163k)

Metric	Value
Accuracy	0.9922
Precision (macro)	0.9931
Recall (macro)	0.9910
F1-score (macro)	0.9920

The MLP successfully captured the MI–SF phase boundary with >99% accuracy across the entire domain. The decision surface remained smooth and consistent with the expected lobe structure of the Bose–Hubbard model.

(b) Reduced Dataset – 30k

Metric	Value
Accuracy	0.9934
F1-score	0.9934
Precision	0.9934
Recall	0.9935

Even with 30k samples (only ~18% of full data), the network preserved near-perfect reconstruction of the phase boundary, demonstrating **strong robustness to data reduction**.

(c) Reduced Dataset – 10k

Metric	Value
Accuracy	0.9821
F1-score	0.9821

Boundary smoothing slightly degraded but overall structure remained faithful. Minor over-smoothing in extrapolated regions was observed, particularly near higher $\bar{U}$ values.

(d) Reduced Dataset – 6k

Metric	Value
Accuracy	0.9813
F1-score	0.9810

With just ~4% of original data, performance still exceeded 97%. The phase lobes were less sharply defined, but the **topology of MI and SF regions** remained intact — a sign of strong generalization given such limited data.

Visualization Summary

- The **decision boundary plots** clearly reproduced the Mott lobes typical of the Bose–Hubbard phase diagram.
- In extrapolated grids, MLP correctly predicted insulating and superfluid domains beyond the training region ($\bar{U} > 25, \bar{\mu} > 60$).
- **Confusion matrices** showed highly symmetric classification — minimal bias toward either phase.

Conclusion

The MLP model provides an **excellent data-driven baseline** for phase classification:

- High accuracy (>99%) even with moderate downsampling.
- Smooth and physically consistent phase boundaries.
- Minimal overfitting and strong extrapolation behavior.

However, since it lacks explicit physics constraints, it serves as a **benchmark** rather than a physics-aware model.

This motivates the transition to **Physics-Informed Neural Networks (PINNs)**, which incorporate physical priors and improve reliability in sparsely sampled or extrapolated regimes.

KERAS BASED NN

🧠 Full-Length Description: Keras Neural Network Model

The **Keras Neural Network (Keras-NN)** model was developed to classify the **quantum phases** (Mott Insulator vs. Superfluid) of the **Bose–Hubbard model** using the two primary system parameters — the normalized interaction strength $\bar{U}$ and chemical potential $\bar{\mu}$. The dataset used for training and evaluation was derived from the numerically simulated phase data file `phase_t1.0_U25.0-(0.1)_Mu60.0-(0.1).csv`, containing over **160,000 samples**. Each data point consists of $(\bar{U}, \bar{\mu}, \rho, \phi^2)$, where ρ is the density and ϕ^2 the superfluid order parameter. The **phase label** was constructed physically:

$$\text{Phase} = \begin{cases} 0 & \text{if } \rho \approx \text{integer and } \phi^2 \approx 0 \text{ (Mott Insulator)} \\ 1 & \text{otherwise (Superfluid)} \end{cases}$$

Before training, all features were standardized using `StandardScaler` to ensure stable optimization. The neural network was implemented using the **Keras Sequential API** with the following architecture:

- **Input layer:** 2 neurons $(\bar{U}, \bar{\mu})$
- **Hidden Layer 1:** 128 neurons, ReLU activation
- **Dropout Layer:** 20% dropout (for regularization)
- **Hidden Layer 2:** 128 neurons, ReLU activation
- **Output Layer:** 1 neuron, Sigmoid activation

The network was trained with **Adam optimizer (learning rate = 0.001)** and **binary cross-entropy loss**, for **50 epochs**, with a **batch size of 32**. The model outputs a probability between 0 and 1 representing the likelihood of the Superfluid phase, which is thresholded at 0.5 to determine the binary class.

⚙️ Dataset Reduction & Testing Setup

To examine scalability and data efficiency, reduced versions of the dataset were also constructed using a **sampling strategy** that preserves deep Mott, deep Superfluid, and intermediate (boundary) regions. The typical splits were in ratios such as:

- **Full dataset:** ~163,000 samples
- **Reduced datasets:** ~30k, ~10k, ~6k, ~4k, and ~1k samples

Each reduced dataset maintained phase balance using stratified sampling.

All models were trained on **75% of the data** and evaluated on a **25% held-out test set**.

📊 Results

The model achieved excellent classification performance, nearly matching the MLP baseline and slightly outperforming it in stability across data reductions.

Dataset Size	Accuracy	F1-score	Notes
Full (~160k)	0.9900	0.9896	Very high accuracy, smooth decision boundary
30k	0.9812	0.9812	Minor drop, still excellent
10k	0.9528	0.9526	Phase boundaries start softening
6k	0.9467	0.9463	Still retains clear MI-SF distinction
4k	0.9214	0.9200	Overfitting evident, yet captures trend qualitatively

Across all scales, **precision and recall remain balanced** (~0.98–0.99 for full data, ~0.95 for reduced), showing no bias toward either phase.

The confusion matrix confirms minimal phase misclassification, particularly at boundary regions where ϕ^2 transitions continuously.

The **predicted phase diagram** reproduces the expected Mott lobes and the Superfluid envelope accurately. Extrapolation to unseen regions ($\bar{U} > 25$, $\bar{\mu} > 60$) remained qualitatively consistent, though with minor overprediction of the SF domain—an expected behavior in purely data-driven models without physical regularization.

🌱 Detailed Conclusion

The Keras neural network serves as a **modernized deep-learning version** of the classical MLP classifier, offering smoother interpolation, better generalization under large data, and enhanced robustness with dropout regularization.

However, despite its strong accuracy within the training regime, the model lacks built-in physical constraints—leading to **mild inaccuracies under extrapolation**.

This observation directly motivated the integration of **physics-informed neural network (PINN)** losses and **Bayesian uncertainty** in later models to improve interpretability and predictive reliability.

📄 Poster Summary (Compact Version)

Model: Keras Neural Network (Deep MLP, $2 \rightarrow 128 \rightarrow 128 \rightarrow 1$, ReLU + Dropout 0.2)

Dataset: Full Bose–Hubbard simulation (163k points, features $\bar{U}, \bar{\mu}$)

Optimizer / Loss: Adam (lr=0.001), Binary Cross-Entropy

Accuracy: 0.99 (Full), 0.98 (30k), 0.95 (10k), 0.92 (4k)

Key Observation: Accurately captures Mott–Superfluid boundaries within training domain; minor extrapolation drift due to lack of physics regularization.

Conclusion: Excellent baseline deep learning model; establishes reference for introducing physics-informed and Bayesian enhancements.

PINN+KERAS NN

Full Model Description: PINN + Keras Hybrid

Dataset and Setup

The model is trained on the Bose–Hubbard phase dataset `phase_t1.0_U25.0-(0.1)_Mu60.0-(0.1).csv`, which spans a dense grid of interaction strength $\bar{U} \in [0, 25]$ and normalized chemical potential $\bar{\mu} \in [0, 60]$. Each sample contains the condensate fraction ϕ^2 , local density ρ , and derived phase label (0 = Mott Insulator, 1 = Superfluid).

The dataset undergoes optional **physics-aware sampling** to construct reduced subsets emphasizing boundary regions between phases—critical for training stability and generalization. Subset sizes range from 4,684 to 163,401 samples.

The input features are $(\bar{U}, \bar{\mu})$, and the supervised outputs are multi-targeted:

- **Phase** (binary classification, via sigmoid),
- **Condensate fraction ϕ^2** and **density ρ** (regression targets),
- **Physics residual** (consistency constraint derived via the surrogate model).

All inputs are normalized and randomly split into training (75%) and validation (25%) sets.

Surrogate Model Construction

Before defining the PINN, a **frozen surrogate model** is trained using a 2-input $(\bar{U}, \bar{\mu}) \rightarrow 2$ -output (ϕ^2, ρ) regressor:

- **Architecture:** $2 \rightarrow 128 \rightarrow 64 \rightarrow 2$ dense layers (ReLU activations).
- **Loss:** Mean Squared Error.
- **Optimizer:** Adam ($\text{lr} = 1 \times 10^{-3}$).
- **Epochs:** 200, batch size = 512, validation split = 0.15.

The surrogate captures the smooth nonlinear mapping of condensate fraction and density, yielding mean absolute errors below 10^{-3} .

Once trained, this model is frozen and re-used within the physics loss term of the hybrid PINN.

PINN + Keras Model Architecture

The hybrid PINN network extends the Keras backbone by integrating **physics residuals** from the surrogate model as an auxiliary output channel:

- **Inputs:** 2D vector $(\bar{U}, \bar{\mu})$.
- **Hidden Layers:** (128, 64) neurons with `tanh` activations (for smoother physical representation).
- **Outputs:**
 1. **phase:** sigmoid activation (binary classification).
 2. ϕ^2 : ReLU (regression).
 3. ρ : ReLU (regression).
 4. **phys_res:** physics residual computed as $\phi^2 - \phi^2_{\text{surrogate}}$ to enforce data–physics agreement.
- **Loss Composition:**

$$\mathcal{L} = \lambda_{\text{phase}} \mathcal{L}_{\text{BCE}} + \mathcal{L}_{\phi^2} + \mathcal{L}_{\rho} + \lambda_{\text{phys}} \|\phi^2 - \phi^2_{\text{surrogate}}\|^2$$

where λ_{phys} is **ramped** during training from $0 \rightarrow 1$ between epochs 5 and 40 via a `LambdaRamp` callback.

- **Optimizer:** Adam ($\text{lr} = 1 \times 10^{-3}$), ReduceLROnPlateau scheduler, batch size = 256.
- **Epochs:** 200 (full dataset); early stopping optional.

Results and Observations

Across all dataset scales, the PINN + Keras hybrid achieves outstanding classification accuracy and physical consistency:

Dataset Size	Accuracy	F1-Score	Comments
163,401	0.9988	0.9988	Excellent global fit, near-perfect phase boundary reproduction.
30,684 (reduced)	0.9909	0.9909	Boundary-aware sampling maintains >99 % accuracy.
10,684	0.9613	0.9613	Smooth transition curve still captured; minor blurring at edges.
6,684	0.9378	0.9379	Accuracy slightly drops near critical $\bar{\mu}$ range.
4,684	0.9441	0.9442	Surrogate physics retains stability even with ~70 % data reduction.

The predicted phase diagrams closely match the true quantum phase boundaries, and the ϕ^2 and ρ fields show physically meaningful gradients across $\bar{U}$ – $\bar{\mu}$ space. Error maps confirm strong alignment with ground-truth condensate and density profiles, validating the surrogate’s corrective influence.

Detailed Conclusion

The **PINN + Keras hybrid** successfully merges data-driven learning with embedded physical priors. Unlike the plain MLP or pure Keras NN, this model constrains its latent representations through the physics-based surrogate, ensuring extrapolative reliability even beyond the training domain. The adaptive λ_{phys} ramping mechanism prevents early-epoch instability while gradually enforcing physical consistency.

Overall, this architecture balances **data fidelity**, **physics regularization**, and **training stability**, achieving both interpretability and precision across full and reduced datasets. It demonstrates a scalable template for physics-guided neural inference in quantum phase prediction.

Poster-Ready Summary

Model Description (PINN + Keras Hybrid):

The Physics-Informed Keras Neural Network integrates a data-driven classifier with a frozen surrogate enforcing physical consistency between predicted condensate fraction (φ^2) and density (ρ). Built with two `tanh` hidden layers (128, 64) and optimized via Adam ($\text{lr} = 1\text{e-}3$), it introduces a dynamic physics loss term ramped from $0 \rightarrow 1$ during training.

Results:

- Accuracy $\approx 99.9\%$ on full data, $\approx 94\%$ on reduced sets.
- Retains accurate MI-SF boundary reconstruction even under 70 % data reduction.
- Physics-guided loss ensures smooth, physically consistent φ^2 and ρ fields.

Conclusion:

The hybrid model surpasses purely data-driven baselines by embedding surrogate-based physical constraints, achieving high fidelity and robust generalization across data scales—an essential step toward interpretable, physics-aware learning of quantum phase transitions.

PINN+Bayesian (MC dropout)

Detailed Model Description

The **Physics-Informed Neural Network (PINN) with a Bayesian surrogate** integrates data-driven learning with physics constraints to accurately predict quantum phases, condensate fraction ϕ^2 , and density ρ in the Bose–Hubbard model. The core idea is to use a **pretrained Bayesian surrogate neural network** as a differentiable physics-informed constraint while simultaneously training the main PINN to classify phases and predict observables. This approach leverages both experimental/computational data and known physical relationships encoded in the surrogate, improving extrapolation reliability beyond the training domain.

Bayesian Surrogate:

The Bayesian surrogate is a fully connected Keras neural network with two hidden layers (128 and 64 neurons) and ReLU activations, trained on the full Bose–Hubbard dataset ($X = [\bar{U}, \bar{\mu}], Y = [\phi^2, \rho]$) for 200 epochs with batch size 512 using the Adam optimizer (learning rate 10^{-3}) and mean-squared error (MSE) loss. Dropout (rate = 0.2) is applied **during both training and inference** to model predictive uncertainty via Monte Carlo sampling. After training, the surrogate is frozen and used to provide a differentiable physics residual in the main PINN.

PINN Architecture:

The main PINN accepts the normalized parameters $[\bar{U}, \bar{\mu}]$ as inputs and contains two hidden layers (128 and 64 neurons) with `tanh` activations. It outputs four branches:

1. **Phase:** Binary classification using sigmoid activation, optimized via binary cross-entropy.
2. **Condensate fraction ϕ^2 :** Regression output, optimized via MSE.
3. **Density ρ :** Regression output, optimized via MSE.
4. **Physics residual:** Difference between the PINN-predicted ϕ^2 and Bayesian surrogate output, scaled by a trainable physics weight λ_{phys} .

The **total loss** is a weighted sum:

$$\mathcal{L}_{\text{total}} = \lambda_{\text{phase}} \mathcal{L}_{\text{BCE}}(\text{phase}) + \mathcal{L}_{\text{MSE}}(\phi^2) + \mathcal{L}_{\text{MSE}}(\rho) + \lambda_{\text{phys}} \underbrace{\frac{1}{N} \sum_i (\phi_i - \phi_i^{\text{surrogate}})^2}_{\text{physics residual}}$$

To ensure stable incorporation of physics information, λ_{phys} is gradually ramped from 0 to 1 over the first 40 epochs using a **custom callback**. The model is trained for 200 epochs with batch size 256 using Adam (learning rate 10^{-3}), with 15% of the data reserved for validation.

Datasets:

- **Full dataset:** 164k points, covering the entire $[\bar{U}, \bar{\mu}]$ parameter space with labeled phases (Mott = 0, Superfluid = 1) using a physics-informed labeling function.
- **Reduced datasets:** Smart sampling selects deep Mott ($\phi^2 \approx 0$, integer density), deep Superfluid ($\phi^2 > 10^{-2}$), and boundary regions with ambiguous phases to create subsets of ~30k, 10k, 6k, and 4k points. This tests model performance under limited data scenarios while focusing on critical phase boundary regions.

Uncertainty Quantification:

The Bayesian surrogate enables Monte Carlo dropout predictions to estimate uncertainty in the PINN outputs. 50 forward passes are performed with dropout active, providing mean and standard deviation of phase, ϕ^2 , and ρ predictions across the parameter grid.

Results

Full dataset (164k points):

- Phase classification accuracy: 99.33%
- Weighted F1-score: 0.9933
- Precision/recall for Mott: 0.9951 / 0.9933, for Superfluid: 0.9908 / 0.9932

Reduced datasets:

Dataset	Accuracy	Mott F1	SF F1	Weighted F1
30k	98.79%	0.9895	0.9857	0.9879
10k	97.67%	0.9797	0.9726	0.9766
6k	97.41%	0.9771	0.9704	0.9741
4k	96.67%	0.9703	0.9622	0.9667

- Predictions of ϕ^2 and ρ closely match the ground truth, with small error distributions concentrated around the phase boundary.
- Monte Carlo uncertainty maps highlight ambiguous boundary regions, providing valuable insight into model confidence.
- Even with aggressive dataset reduction, the PINN maintains high accuracy and robust phase boundary identification, outperforming purely data-driven MLP/Keras NN baselines.

Detailed Conclusion

The **PINN + Bayesian surrogate model** demonstrates that embedding physical knowledge into neural networks significantly improves predictive reliability and generalization in quantum many-body systems. By combining a Bayesian surrogate with physics-informed loss terms, the network not only achieves near-perfect phase classification (>99% on full dataset) but also provides uncertainty estimates essential for trust in extrapolated regions. Reduced datasets highlight the model's ability to maintain strong performance with fewer data points, thanks to targeted sampling and physics constraints. The physics residual ensures that the predicted condensate fraction aligns with the surrogate-based physical model, preventing unphysical extrapolations typical of purely data-driven approaches. Overall, this approach enables **accurate, interpretable, and uncertainty-aware predictions of Bose–Hubbard phase diagrams**.

Poster-Ready Summary

Model Description:

A Physics-Informed Neural Network integrates a pretrained Bayesian surrogate to enforce physical consistency while predicting phase, ϕ^2 , and ρ from Bose–Hubbard parameters. The network uses two hidden layers (128,64 neurons, $\tanh$) and outputs a physics residual from the surrogate. Loss combines binary cross-entropy, MSE, and the physics residual, with ramped weighting for stable training.

Results:

- Full dataset: 99.33% accuracy, $F1 > 0.993$.
- Reduced datasets: 96–98% accuracy depending on size.
- Physics residual ensures physically consistent predictions.
- Monte Carlo dropout provides uncertainty estimates highlighting ambiguous boundaries.

Conclusion:

The PINN + Bayesian surrogate combines data-driven learning with physics knowledge to deliver accurate, interpretable, and uncertainty-aware predictions of Bose–Hubbard quantum phases, outperforming traditional MLP/Keras models, especially in extrapolation or reduced-data regimes.